






188: End To End Project With PDF Query RAG With Langchain And AstraDB





Code is self explanatory

+

-

DataStax Astra DB Serverless Documentation

Overview >

Getting Started >

Vector Search ✓

Quickstart

Examples

Query vector data with CQL >

Working with embeddings

Indexing

About Vector Search

Planning >

Securing >

DATASTAX ASTRA DB SERVERLESS DOCUMENTATION > VECTOR SEARCH

Vector Search

Vector Search enhances machine learning models by allowing similarity comparisons of embeddings, which are mathematical representations of high dimensional data.

As a capability of Astra DB, Vector Search supports various Large Language Models (LLM). Since these LLMs are stateless, they rely on a vector database like Astra DB to store their embeddings. You can expedite your vector-based similarity searches by using **Serverless Cassandra with Vector Search**, making it easier to develop your LLM-powered applications.

Here are some of the top reasons to use our vector search:

- Discover connections by comparing similarities within vast collections of documents
- Match product recommendations, find similar images, and more
- Seamless integration with LLMs
- Scale effortlessly to handle large volumes of embeddings

What is Apache Cassandra®?

Apache Cassandra is an open source NoSQL distributed database trusted by thousands of companies for scalability and high availability without compromising performance. Linear scalability and proven fault-tolerance on commodity hardware or cloud infrastructure make it the perfect platform for mission-critical data.



Hybrid

Masterless architecture and low latency means Cassandra will withstand an entire data center outage with no data loss—across public or private clouds and on-premises.




Fault Tolerant

Cassandra's support for replicating across multiple datacenters is best-in-class, providing lower latency for your users and the peace of mind of knowing that you can survive regional outages. Failed nodes can be replaced with no downtime.



Focus on Quality

To ensure reliability and stability, Cassandra is tested on clusters as large as 1,000 nodes and with hundreds of real-world use cases and schemas tested with replay, fuzz, property-based, fault-injection, and performance tests.



DATASTAX

krishnaik06@gmail.com

Home

Databases

- langchain_db

Streaming

Billing

Tokens

Settings

Integrations

Sample Apps

Documentation

Command Palette

Free Plan Upgrade

Accelerate on Astra with 1,000+ Credits

Welcome to Astra, Krish 🙌

Accelerate your workflows, access recently visited resources, and explore Astra's integrations and documentation!

PREVIEW

Introducing a New Vector Developer Experience

Try our new, intuitive vector portal experience and schemaless API for faster Gen AI and RAG (retrieval-augmented generation) app development.

Enable Preview

Guides & Examples

Q&A Search with LangChain 🚀

Perform a text similarity search on HuggingFace datasets using Astra DB, LangChain and CassIO.

View Quickstart

Astra DB with OpenAI's Assistants API 🧙

Change one line of code and leverage Astra DB as your vector store with OpenAI's new Assistants API.

View Video

Vector Search using the JSON API 📄

Perform a text similarity search against a movie dataset using Astra DB and the Mongoose powered JSON API for JavaScript.

View Example

Quick Access

Create a Database

Generate Token

Invite your team

Recent Resources

Recommended Integrations

