

LAMINI

Lamini Docs

Lamini Docs

Introduction

Overview

About

Authenticate

Enterprise Install

FAQ

Inference

Quick Start

Prompt Tuning

JSON Output

Batching

Retrieval Augmented Generation (RAG)

Tuning

Quick Start

Memory Tuning

Hyperparameters

Large Data Files

Evaluation

Dashboard

Playground

Applications

Pipeline

Overview

Welcome to Lamini

Lamini is an integrated LLM inference and tuning platform, so you can seamlessly tune models that achieve exceptional factual accuracy while minimizing latency and cost. Lamini Enterprise runs in your environment - even air-gapped - or you can use our GPUs in Lamini Cloud.

Goal	Go to
2 steps to start using LLMs on Lamini Cloud	Quick Start
95% accuracy and beyond	Memory Tuning
LLM inference that's 100% guaranteed to match your schema	JSON Output
High throughput inference (52x faster)	Iteration Batching
What makes Lamini unique?	About
Use cases and recipes	Examples

Table of contents

Welcome to Lamini

Revisit the video again